

The μ Specification — r0.3

Preface

μ (stylized μ) is an extremely simple programming language. This document contains a complete specification of the language followed by a series of appendices elucidating motivation, design decisions, implementation strategies, and so on. These appendices are not part of the formal language specification and in the event of a conflict between the specification and the content of an appendix the specification is authoritative.

This is revision r0.3 of the specification.

1 Source

1.1 Encoding

A μ program (recommended extensions μ , $\mu\mu$) consists of a series of bytes. The behaviors of bytes with the high bit set is implementation defined. As a consequence a conforming implementation MAY choose to interpret μ code as text in ASCII, UTF-8, or another similar encoding. This document will adopt such an interpretation when displaying μ source.

1.2 Syntax

1.2.1 Grammar

```
<program> ::= <ws> <expr> <ws>

<expr> ::= <atom> | <list>

<atom> ::= <ordinary> { <ordinary> }

<list> ::= <lparen>
          { <ws> <expr> }
          [ <ws> <dot> <ws> <expr> ] <ws>
          <rparen>

<ws> ::= { <tab> | <lf> | <cr> | <sp> }

<lparen> ::= 0x28 ;; LEFT PARENTHESIS '('
<rparen> ::= 0x29 ;; RIGHT PARENTHESIS ')'
<dot> ::= 0x2e ;; FULL STOP '.'

<number> ::= 0x23 ;; NUMBER SIGN '#'
<quote> ::= 0x22 ;; QUOTATION MARK '"'
```

```
<tab> ::= 0x09 ;; HORIZONTAL TAB '\t'
<lf> ::= 0x0a ;; LINE FEED '\n'
<cr> ::= 0x0d ;; CARRIAGE RETURN '\r'
<sp> ::= 0x20 ;; SPACE ' '

<ordinary> ::= <any> - <special>

<any> ::= 0x20-0x7e ;; ALL PRINTABLE ASCII
<special> ::= <lparen> | <rparen> | <dot>
              | <number> | <quote>
              | <tab> | <lf> | <cr> | <sp>
```

The characters # and " are reserved.

1.2.2 Interpretation

Each distinct μ atom in the source is an *atom name*, a sequence of bytes, and is given a distinct positive *atom number*.^{[1][2]} How different byte sequences are mapped to numbers is completely implementation defined with the only restrictions being that:

- Atom numbers are between 1 and $2^{31} - 1$.
- Two atoms with the same name within the same execution of a program get the same atom number.
- Two atoms with different names within the same execution of a program get different atom numbers.

For example, all instances of the atom with name hello (hex 68 65 6c 6c 6f) are to be assigned the same atom number (say, 3) which must be different from the atom number for world (hex 77 6f 72 6c 64) (say, 7) but what exactly the atom numbers are can vary between implementations, programs, or even different executions of the same program at the implementer's discretion.

An implementation MAY assume that there are no more than $2^{11} - 1$ distinct atom names referred to in a given source text, that each atom name has a length of at most $2^7 - 1$ bytes, and that the combined length of these names (excluding duplicates) is no more than $2^{15} - 2^{11}$ bytes.^[3]

^[1]See 2.1 Atoms.

^[2]See TODO B Implementation Strategies.

^[3]These restrictions permit an implementation to allocate its atom names into a singular region of memory of size i32::MAX separated by NULLs, for instance.

A list in the source code is constructed of cons cells^[4] according to the following algorithm:

```
def parse_list(tokens):
    eat a lparen from tokens
    list = parse_list_inner(tokens)
    eat a rparen from tokens
    return list

def parse_list_inner(tokens):
    if ( the next token in tokens is dot ):
        eat a dot from tokens
        return parse an expression from tokens
    elif ( the next token in tokens is rparen ):
        return the () atom
    else:
        head = parse an expression from tokens
        tail = parse_list_inner(tokens)
        return a new cons cell of head and tail
```

A few things to note about this algorithm:

- `()` becomes the `()` atom. For this reason the `()` atom is also called `nil` and is usually denoted `()`.
- A list like $(i_1\ i_2\ i_3)$ denotes an ordinary linked list $i_1 \rightarrow i_2 \rightarrow i_3 \rightarrow \mathbf{NIL}$.
- A dotted pair like $(l\ .\ r)$ denotes a simple cons cell (l, r) .
- A dotted list like $(i_1\ i_2\ i_3\ .\ r)$ denotes a modified linked list (with non-nil tail) $i_1 \rightarrow i_2 \rightarrow i_3 \rightarrow r$.
- An ordinary list like $(i_1\ i_2\ i_3\ \dots\ i_n)$ is equivalent to the dotted $(i_1\ .\ (i_2\ .\ (i_3\ .\ (\dots\ (i_n\ .\ ())))))$.

2 Types

`μ` is dynamically typed with exactly two types: atoms and cons cells.

2.1 Atoms

An atom is fundamentally an *unsigned 31 bit integer*.^[5]

The `()` atom represents the empty list in list-related contexts.

Some positive atoms correspond to particular byte sequences (atom names) in the source code.^[6] Not all atom values have a corresponding name nor is there a standard mechanism for converting between atoms and names during the execution of a program. This correspondence is simply a mechanism to allow source code to provide human-readable names to arbitrary symbols rather than a property of the language's runtime.

^[4]See 2.2 Cons Cells.

^[5]See B Implementation Strategies for a rationale for this unusual choice.

^[6]See 1.2.2 Interpretation.

2.2 Cons Cells

A cons cell is simply an ordered pair of items (each of which may be either an atom or another cons cell).

The first item in the pair is called the *head* and the second is called the *tail*.

A list refers to a particular structure formed of either the `()` atom or a cons cell whose tail is itself a list.^[7]

3 Environments

At the core of `μ`'s scoping rules is the environment, a mapping from atoms to arbitrary values. This mapping is defined as a list of pairs of positive atoms and values. The first element in each pair is a variable to bind and the second is an value to bind to that variable.

For example:

```
( (hello . hello)
  (sometest . (a b c))
  (a_variable . a_value) )
```

is an environment mapping the atom `hello` to itself, the atom `sometest` to the list `(a b c)`, and the atom `a_variable` to the atom `a_value`.

An environment may contain multiple bindings for the same atom in which case the earlier binding *shadows* the later one, effectively overriding it^[8].

`()` is a legal environment containing no mappings.

3.1 Lookup

Looking up the value corresponding to an atom in an environment is defined according to the following algorithm:

```
def lookup(sym, env):
    if ( env is empty ):
        return sym
    else:
        binding = the head of env
        if ( the head of binding is sym):
            return the tail of binding
        else:
            return lookup(sym, the tail of env)
```

Note that duplicate bindings are resolved in favor of the first, and missing bindings resolve to themselves.

^[7]See 1.2.2 Interpretation for how these are denoted.

^[8]The shadowed binding is still accessible via `~env`, though.

4 Pattern Matching

A pattern is a value which can be matched against another value (called the scrutinee) in some base environment to produce a new environment which contains everything in the base environment in addition to new bindings representing components of the value.

Every value which does not contain duplicate positive atoms is a valid pattern.

Matching a pattern against a value follows the following rules:

- `()` matches the value `()` and introduces no bindings. Attempting to match this pattern against any other object is undefined behavior.^[9]
- A positive atom matches any value and introduces a binding from the pattern atom to the value.
- A cons cell matches a cons cell value, recursively pattern matching the value's head and tail against its head and tail. The order in which the bindings from the head and tail are included in the final environment is implementation defined and may even be non-deterministic. Attempting to match a cons cell pattern against any other object is undefined behavior.
- Introduced bindings shadow bindings in the base environment (i.e. are placed before) but do not remove them. This allows careful environment manipulation to recover the base environment which is necessary for writing fully hygienic macros.

4.1 Examples

- Matching the expression `()` against the pattern `()` in the base environment `((a . b))` produces the environment `((a . b))`.
- Matching the expression `(a b c d)` against the pattern `(x y . z)` in the base environment `()` produces some permutation of the environment `((x . a) (y . b) (z . (c d)))`
- Matching the expression `(a b c)` against the pattern `(x y)` is undefined behavior.
- Matching the expression `(a b)` against the pattern `(x y)` in the base environment `((x . xx))` may produce either the environment `((x . a) (y . b) (x . xx))` or `((y . b) (x . a) (x . xx))` but MUST NOT produce the environment `((x . a) (y . b))`.

^[9]See 7 Undefined Behavior.

5 Evaluation

The heart of the language is evaluation. In fact, executing a program simply consists of parsing it then evaluating it in an empty environment.

An *expression* is any value to be evaluated.

Evaluation takes place in an environment e .^[10]

Evaluating an expression behaves differently depending on whether the expression is the 0 atom, a positive atom, or a non-empty list.

5.1 Zero Atom

The 0 atom `()` evaluates to itself.

5.2 Positive Atoms

A positive atom a evaluates to `lookup(a , e)`.^[11]

5.3 Lists

If a cons cell is evaluated it should be a valid list. An implementation MAY reject attempting to evaluate a non-list cons cell or MAY try to interpret such an expression cohesively.

A list is evaluated as a call by first evaluating its head and then, based on the result (referred to as the *receiver*), performing an operation on the remaining elements.

We will use the meta syntax of angle brackets $\langle expr \rangle_{env}$ to denote an expression which evaluates to $expr$ in the environment env and the syntax of an arrow with environment above (\xrightarrow{env}) to denote the evaluation of an expression to a value in an environment.

5.3.1 Builtins

If the receiver is an atom it should belong to the following list of atoms with builtin behavior. Notice that with the exception of `()` which has no corresponding name these builtin receivers all have names starting `~` (hex 7e 7e).

`()` Called "quote", but represented by the 0 atom for technical reasons.^[12] Evaluates to its first argument *unevaluated*.

$$(\langle () \rangle_e i) \xrightarrow{e} i$$

^[10]See 3 Environments.

^[11]See 3.1 Lookup for the definition of `lookup()`.

^[12]Using the 0 atom as the quote builtin macro ensures that the quote macro is always accessible as the 0 atom is the only atom which cannot be rebound. This, in turn, ensures that any value is always accessible by invoking quote. The ability to access arbitrary content reliably regardless of environment is critical for hygienic macros to be possible.

~true	Takes two arguments and tail evaluates ^[13] to its first argument. DOES NOT evaluate its second argument.	~sub	Takes two arguments and evaluates them. They should both be atoms. Returns the difference of the two atoms modulo 2^{31} .
	$((\sim\text{true})_e \langle t \rangle_e f)_e \xrightarrow{e} t$		$((\sim\text{sub})_e \langle a \rangle_e \langle b \rangle_e) \xrightarrow{e} (a - b) \bmod 2^{31}$
~false	Takes two arguments and tail evaluates to its second argument. DOES NOT evaluate its first argument.	~and	Takes two arguments and evaluates them. They should both be atoms. Returns the bitwise conjunction (bitwise and) of the two atoms
	$((\sim\text{false})_e t e \langle f \rangle_e) \xrightarrow{e} f$		$((\sim\text{and})_e \langle a \rangle_e \langle b \rangle_e) \xrightarrow{e} (a \wedge b)$
~head	Takes an argument and evaluates it. The result should be a cons cell. Returns the head of that cons cell.	~or	Takes two arguments and evaluates them. They should both be atoms. Returns the bitwise disjunction (bitwise or) of the two atoms
	$((\sim\text{head})_e \langle (h . t) \rangle_e) \xrightarrow{e} h$		$((\sim\text{or})_e \langle a \rangle_e \langle b \rangle_e) \xrightarrow{e} (a \vee b)$
~tail	Takes an argument and evaluates it. The result should be a cons cell. Returns the tail of that cons cell.	~not	Takes one argument and evaluates it. It should be an atom. Returns the bitwise negation (bitwise not) of the atom
	$((\sim\text{tail})_e \langle (h . t) \rangle_e) \xrightarrow{e} t$		$((\sim\text{not})_e \langle a \rangle_e) \xrightarrow{e} (\neg a)$
~cons	Takes two arguments, evaluates them, and returns a new cons cell constructed from the first and second arguments.	~shl	Takes two arguments and evaluates them. They should both be atoms. The second argument should be less than 31. Returns the first shifted left by the second number of bits.
	$((\sim\text{cons})_e \langle h \rangle_e \langle t \rangle_e) \xrightarrow{e} (h . t)$		$((\sim\text{shl})_e \langle a \rangle_e \langle b \rangle_e) \xrightarrow{e} (a << b)$
~lte	Takes two arguments and evaluates them. Returns ~true if the first is less than or equal to the second and ~false otherwise. Atoms are compared according to their numbers. Cons cells are always less than atoms. Two distinguishable cons cells do not compare equal. Two indistinguishable cons cells may or may not compare equal.	~shr	Takes two arguments and evaluates them. They should both be atoms. The second argument should be less than 31. Returns the first shifted right by the second number of bits.
	$((\sim\text{lte})_e \langle a \rangle_e \langle b \rangle_e) \xrightarrow{e} \begin{cases} \sim\text{true} & \text{if } a \leq b \\ \sim\text{false} & \text{otherwise} \end{cases}$		$((\sim\text{shr})_e \langle a \rangle_e \langle b \rangle_e) \xrightarrow{e} (a >> b)$
~eq	Takes two arguments, evaluates them, and returns ~true if the first is equal to the second and ~false otherwise. Atoms are compared according to their numbers. Cons cells are always less than atoms. Two distinguishable cons cells do not compare equal. Two indistinguishable cons cells may or may not compare equal.	~env	Takes no arguments. Returns the environment.
	$((\sim\text{eq})_e \langle a \rangle_e \langle b \rangle_e) \xrightarrow{e} \begin{cases} \sim\text{true} & \text{if } a = b \\ \sim\text{false} & \text{otherwise} \end{cases}$		$((\sim\text{env})_e) \xrightarrow{e} e$
~add	Takes two arguments and evaluates them. They should both be atoms. Returns the sum of the two atoms modulo 2^{31} .	~sys	Takes one argument and evaluates it. Usually returns an implementation defined receiver based on its argument. See 5.3.3 The ~sys Builtin for a complete definition.
	$((\sim\text{add})_e \langle a \rangle_e \langle b \rangle_e) \xrightarrow{e} (a + b) \bmod 2^{31}$		$((\sim\text{sys})_e \text{sym}) \dots \xrightarrow{e} \text{See 5.3.3}$

^[13]See 6 Required Optimizations.

^[14]See 7 Undefined Behavior.

Each builtin takes a fixed number of arguments. Providing the wrong number of arguments to a builtin is undefined behavior.^[14]

Attempting to use any atom that is not on this list of receivers is undefined behavior as is passing an argument of (or that evaluates to, where relevant) a type other than that expected.

5.3.2 User Defined Receivers

If the receiver is a list then it represents a user-defined macro or function and should have exactly 2 or 3 elements, respectively.

The first of these elements is a pattern^[15] and the second is the body. The third element has two purposes:

1. Its presence indicates that this receiver is a function rather than a macro.
2. It defines an environment to use when expanding the function body. Combined with the `~env` builtin, this allows lexical scoping of functions, scoped evaluation as a derived object, and more.^[16]

To evaluate a call involving a user-defined receiver preform the following steps:

1. If the receiver is a function, evaluate all of its arguments, otherwise, leave them unevaluated.
2. Match the list of arguments against the included pattern to produce a new environment.^[17]
 - If the receiver is a function, use the included environment as the base when constructing the new environment.
 - Otherwise, use the calling environment as the base when constructing the new environment.
3. Tail evaluate^[18] the included body in the constructed environment.

Attempting to use any cons-cell which does not form a valid list of either exactly two or three elements where the first element is a valid pattern and the third (if present) is a valid environment as a receiver is undefined behavior.^[19]

5.3.3 The `~sys` Builtin

`~sys` is the main point of extensibility built into `μ`.

Calling the `~sys` with the 0 atom, e.g. as `((() ~sys) ())`, should return a list of pairs mapping *system operation names* to *system operation codes*.^[20] This mapping should remain constant throughout a program's lifetime.

The heads of these pairs, the *system operation names*, are atoms which correspond to human readable names of a sort a programmer may write in their source code. The tails,

meanwhile, the *system operation codes*, should also be atoms but which need not be human readable names.

Calling `~sys` with a system operation code should return a receiver which can be called to perform the desired system operation. The implementation may choose to represent this receiver using any `μ` object except one of the builtin atoms or a list of one or two elements. The receivers may be, in general, macro like, performing a more complex transform on their arguments than simple evaluation.

Using both a readable name and a code and delegating the mapping to the user's program this way allows for an extremely simple implementation of `~sys` using a jump table or similar structure while still allowing readable names to be used to refer to these operations.

6 Required Optimizations

Several points in 5 Evaluation specify to tail evaluate some expression. In such cases implementations MUST perform proper tail-call optimization. This means that evaluating an expression which causes a tail evaluation of another expression which itself causes the tail evaluation of another expression and so on should only take up $O(1)$ space for storage of the current evaluation state (e.g. call stack). This optimization is necessary to ensure programmers can write loops without fear of running out of space.

Implementations must not needlessly duplicate cons cells. Acquiring a cons cell from a binding multiple times should not require additional space for each instance.

Implementations must have a garbage collector. Concretely, repeatedly creating and then no longer using cons cells should not cause memory expenditure to grow without bound.

E Tests offers a list of test cases which, in addition to testing correctness of an implementation, test that these characteristics are met appropriately.

7 Undefined Behavior

Several points throughout this specification refer to certain occurrences as undefined behavior. Implementations are free to adopt arbitrary behavior in these instances including but not limited to:

- Extending the specification behavior
- Producing some kind of error
- Producing a nonsense result
- Halting and catching fire

^[15]See 4 Pattern Matching.

^[16]See TODO.

^[17]See 4 Pattern Matching for the definition of this operation.

^[18]See 6 Required Optimizations.

^[19]See 7 Undefined Behavior.

^[20]For a system which does not provide implementation specific functionality through this mechanism just return `()`.

Appendix A

Background

Appendix B

Implementation Strategies

Appendix C

Derivation of Higher-Level Functionality

Appendix D

Reference Implementation (WebAssembly)

This appendix contains a reference implementation of a μ interpreter written in WebAssembly and the Javascript bindings necessary to embed it on the Web or in Deno.

While the full source code of these files is included in this appendix for completeness, a reader interested in acquiring copies can much more easily obtain them from the reference-implementation directory of the μ project's git repository in clean, plain text.

D.1 WebAssembly Text — mu_.wat

```
;; # `mu_.wat` - A reference implementation for the mu_ programming language
;;
;; `mu_.wat` is handcrafted, unfolded, and well commented WebAssembly text
;; program which can serve suitably both as an implementation in its own right
;; and as a template for ports to other assembly and programming languages.
;; See `mu_.mjs` for Javascript bindings
;;
;; This is version `r0.3i1` which implements `r0.3` of the mu_ specification.
;;
;; To compile, run: `wat2wasm --enable-multi-memory --enable-tail-call mu_.wat`
;;
;; ## Exposed Details
;;
;; All items are uniformly represented as 32 bit signed integers, with positive
;; numbers representing atoms in the obvious way and negative numbers
;; representing cons cells as their negative offsets into the cons cell stack.
;;
;; There is a string yard, a simple memory buffer and allocator `syalloc` which
;; the embedder can write strings into as a pre-condition to calling interpreter
;; methods which expect strings.
;;
;; `mu_.wat` exports the following bindings (in order of definition)
;; - `cons` : Construct a cons cell
;; - `head` : Take the head of a cons cell
;; - `tail` : Take the tail of a cons cell
;; - `lookup` : Look up a symbol in an environment
;; - `match` : Match a value against a pattern in a base environment
;; - `stringyard` : A buffer for the embedder to write strings (such as
;;   program source code or atom names) into
;; - `syalloc` : Allocate space in the string yard for a new string
;; - `inter_string` : Take an offset and length (in bytes) into the string
;;   yard and inter the string onto the internment stack,
;;   returning its atom number
;; - `lookup_interred_string` : Given an atom, return the offset and length
;;   into the string yard of that atom's name,
;;   yielding -1, -1 if it is unnamed
;; - `parse` : Parse a string in the string yard to a mu_ value
;; - `eval` : Evaluate an expression in an environment (then gc)
;; - `register_system_operation` : Given an atom and a funcref register a
;;   system operation with that name handled by
;;   that handler
;; - `gc_get_anchor` : Get an anchor for the garbage collector
;;   Usually, you should call this and pass the result as
;;   the third argument to `eval`. Read the comments in the
;;   garbage collection section which offer more detail on
;;   how garbage collection is implemented before using the
;;   the garbage collection mechanism in any other way
;; - `gc_collect` : Run the garbage collector
;;
;; (C) 2025 Brielle Hoff --- Dual licensed under CC BY-NC 4.0 and MIT.
(module
  ;; ----- Cons Cells -----
  ;; Each cell is (head, tail)
  ;; $cons_cells_top points to the last allocated cons cell
  (memory $cons_cells 1)
  (data (memory $cons_cells) (i32.const 0) "\00\00\00\00\00\00\00\00")
  ;; Points to the last element of the stack
  (global $cons_cells_top (mut i32) (i32.const 0))

  ;; Construct a cell

  (func $cons (export "cons")
    (param $head i32) (param $tail i32)
    (result i32)
    (local $top i32)

    ;; Adjust the stack pointer
    global.get $cons_cells_top
    i32.const 8
    i32.add
    local.tee $top
    global.set $cons_cells_top

    ;; Write the head value
    local.get $top
    local.get $head
    i32.store (memory $cons_cells)

    ;; Write the tail value
    local.get $top
    i32.const 4
    i32.add
    local.get $tail
    i32.store (memory $cons_cells)

    ;; Negate the offset into the stack to produce the representation
    i32.const 0
    local.get $top
    i32.sub

  )

  ;; Take the head of a cell
  (func $head (export "head")
    (param $cell i32)
    (result i32)

    ;; Negate the representation to get the offset into the stack
    i32.const 0
    local.get $cell
    i32.sub

    ;; Read the head value
    i32.load (memory $cons_cells)

  )

  ;; Take the tail of a cell
  (func $tail (export "tail")
    (param $cell i32)
    (result i32)

    ;; Negate the representation to get the offset into the stack and
    ;; shift by four to get the tail instead of the head in one step
    i32.const 4
    local.get $cell
    i32.sub

    ;; Read the tail value
    i32.load (memory $cons_cells)

  )

  ;; ----- Primitive Interpreter Operations -----

  ;; Lookup a symbol in an environment
  (func $lookup (export "lookup")
    (param $symbol i32)
    (param $environment i32)
    (result i32)
    (local $binding i32)

    (loop $loop (result i32)
      local.get $environment
      i32.eqz
      (if (result i32)
        (then
          ;; We've exhausted the environment, map a symbol to itself
          local.get $symbol
        )
        (else
          ;; Get the first binding in the environment
          local.get $environment
          call $head
          local.tee $binding

          call $head
          local.get $symbol
        )
      )
    )

  )
```

```

i32.eq
(if (result i32)
  (then
    ;; If it matches, return the bound value
    local.get $binding
    call $tail
  )
  (else
    ;; Otherwise, continue over the remaining bindings
    local.get $environment
    call $tail
    local.set $environment
    br $loop
  )
)
))
)

;; Match a value against a pattern in a base environment
(func $match (export "match")
  (param $value i32)
  (param $pattern i32)
  (param $environment i32)
  (result i32)

  local.get $pattern
  i32.eqz
  (if (result i32)
    (then
      ;; The pattern is (), don't introduce any bindings
      local.get $environment
    )
    (else
      local.get $pattern
      i32.const 0
      i32.gt_s
      (if (result i32)
        (then
          ;; The pattern is a positive atom, add a binding to the
          ;; environment
          local.get $pattern
          local.get $value
          call $cons
          local.get $environment
          call $cons
        )
        (else
          ;; The pattern is a cons cell, recurse

          ;; Match the value head against the pattern head
          local.get $value
          call $head
          local.get $pattern
          call $head

          ;; Match the value tail against the pattern tail
          local.get $value
          call $tail
          local.get $pattern
          call $tail
          local.get $environment

          call $match
          call $match
        )
      )
    )
  )
)

;; ----- String Yard -----
;; The stringyard is a place for the embedding application to place strings
;; it wishes to pass to the interpreter.
;; Call syalloc(size) to acquire size bytes from the yard to place a string
(memory $stringyard (export "stringyard") 1)
(data (memory $stringyard) (i32.const 0)
  "true..~false.."
  "head..~tail..~cons.."
  "lte...~eq..."
  "add...~sub..."
  "and...~or...~not..."
  "shl...~shr..."
  "env...~sys..."
)
(global $stringyard_top (mut i32) (i32.const 128))

;; Allocate space on the string yard to place a string of size bytes
;; mu_ never uses this internally (nor modifies the string yard at all)
;; it only uses it for its initial strings and as a dumping ground for the
;; embedder to place strings into for it to use.
(func (export "syalloc")
  (param $size i32)
  (result i32)
  (local $top i32)

  ;; Adjust the available space top
  ;; This implementation is extremely simple because we don't do any
  ;; sophisticated memory management for strings.
  global.get $stringyard_top
  local.tee $top
  local.get $size
  i32.add
  global.set $stringyard_top

  local.get $top

)

;; Compare two strings in the string yard
(func $str_eq
  (param $a_off i32)
  (param $b_off i32)
  (param $a_len i32)
  (param $b_len i32)
  (result i32)
  (local $a_end i32)

  ;; Pre-check the string lengths
  local.get $a_len
  local.get $b_len
  i32.eq
  (if
    (then
      ;; Convert len to end for simpler iteration
      local.get $a_off
      local.get $a_len
      i32.add
      local.set $a_end
    )

    ;; Loop over the bytes in the strings
    (block $break_loop
      (loop $loop
        ;; If we've reached the end, the strings are equal
        local.get $a_off
        local.get $a_end
        i32.eq
        (if
          (then
            i32.const 1
            return
          )
        )

        ;; Compare bytes, break if unequal
        local.get $a_off
        i32.load8_u (memory $stringyard)
        local.get $b_off
        i32.load8_u (memory $stringyard)
        i32.ne
        br_if $break_loop

        ;; Increment the offsets into the strings
        local.get $a_off
        i32.const 1
        i32.add
        local.set $a_off
        local.get $b_off
        i32.const 1
        i32.add
        local.set $b_off
      )
    )
    br $loop
  )

  ;; The strings are different
  i32.const 0
)

;; String internment
(memory $string_internment_stack 1)
;; [{ offset: i32, len: u16, system_opcode: u16 }]
(data (memory $string_internment_stack) (i32.const 0)
  "\00\00\00\00" "\00\00" "\00\00" ;; <ensure no string at 0>
  "\00\00\00\00" "\06\00" "\00\00" ;; ~true (6)
  "\08\00\00\00" "\07\00" "\00\00" ;; ~false (7)
  "\10\00\00\00" "\06\00" "\00\00" ;; ~head (6)
  "\18\00\00\00" "\06\00" "\00\00" ;; ~tail (6)
  "\20\00\00\00" "\06\00" "\00\00" ;; ~cons (6)
  "\28\00\00\00" "\05\00" "\00\00" ;; ~lte (5)
  "\30\00\00\00" "\04\00" "\00\00" ;; ~eq (4)
  "\38\00\00\00" "\05\00" "\00\00" ;; ~add (5)
  "\40\00\00\00" "\05\00" "\00\00" ;; ~sub (5)
  "\48\00\00\00" "\05\00" "\00\00" ;; ~and (5)
  "\50\00\00\00" "\04\00" "\00\00" ;; ~or (4)
  "\58\00\00\00" "\05\00" "\00\00" ;; ~not (5)
  "\60\00\00\00" "\05\00" "\00\00" ;; ~shl (5)
  "\68\00\00\00" "\05\00" "\00\00" ;; ~shr (5)
  "\70\00\00\00" "\05\00" "\00\00" ;; ~env (5)
  "\78\00\00\00" "\05\00" "\00\00" ;; ~sys (5)
)

;; Points "One past the end" of the stack
(global $string_internment_stack_top (mut i32) (i32.const 136))

;; Inter a string from the stringyard onto the string internment stack
(func $inter_string (export "inter_string")
  (param $off i32)
  (param $len i32)
  (result i32)
  (local $idx i32)

  ;; Initialize idx to zero
  i32.const 0
  local.set $idx

  ;; Loop over the strings in the internment stack
  (block $scan
    (block $break_loop
      (loop $loop
        ;; If we reached the end of the internment stack, we need
        ;; to add a new string
        local.get $idx
        global.get $string_internment_stack_top
        i32.eq
        br_if $break_loop

        ;; Load the idx'th string from the stack and compare to the
        ;; string to be interred, if equal we found a match
        local.get $idx

```

```

        i32.load (memory $string_internment_stack)
        local.get $idx
        i32.const 4
        i32.add
        i32.load16_u (memory $string_internment_stack)
        local.get $off
        local.get $len
        call $str_eq
        br_if $scan

        ;; Increment idx
        local.get $idx
        i32.const 8
        i32.add
        local.set $idx

        br $loop
    )
)

;; We need to add a new string

;; Store the string onto the internment stack
local.get $idx
local.get $off
i32.store (memory $string_internment_stack)
local.get $idx
i32.const 4
i32.add
local.get $len
i32.store16 (memory $string_internment_stack)
local.get $idx
i32.const 6
i32.add
i32.const 0
i32.store16 (memory $string_internment_stack)

;; Adjust the top of the internment stack
local.get $idx
i32.const 8
i32.add
global.set $string_internment_stack_top
)

local.get $idx

;; Flip bit 29 to decrease the odds the numeric representation occurs
;; by chance when doing ordinary calculations
i32.const 0x20_00_00_00
i32.xor
)

;; Lookup an interred string or return -1 -1 if not an interred string atom
;; String atom detection is on a best effort basis as all string atoms have
;; a numeric value it is always possible that a string atom is created by
;; chance via a numeric calculation.
;; String atoms are given unusual numbers to minimize the likelihood of this
;; but it remains possible.
(func (export "lookup_interred_string")
    (param $idx i32)
    (result i32 i32)

    ;; Undo the flip of bit 29
    local.get $idx
    i32.const 0x20_00_00_00
    i32.xor
    local.set $idx

    (block $unnamed
        ;; Check that the index is in bounds
        local.get $idx
        global.get $string_internment_stack_top
        i32.ge_u
        br_if $unnamed

        ;; Check that the index is properly aligned
        local.get $idx
        i32.const 0x7
        i32.and
        i32.const 0
        i32.ne
        br_if $unnamed

        ;; This could be an interred string, load its properties

        ;; Load the offset
        local.get $idx
        i32.load (memory $string_internment_stack)

        ;; Load the length
        local.get $idx
        i32.const 4
        i32.add
        i32.load16_u (memory $string_internment_stack)

        return
    )

    ;; This isn't an interred string, return the failure sentinel
    i32.const -1
    i32.const -1
    return
)

;; ----- Parsing -----

;; Parse mu_source text to an expression
(func $parse (export "parse")
    (param $off i32)
    (param $len i32)

    (result i32)
    (local $end i32)

    ;; Convert len to end for easier iteration
    local.get $off
    local.get $off
    local.get $len
    i32.add
    local.set $end

    ;; Skip leading whitespace
    call $parse_skip_ws

    ;; Parse an expression
    local.get $end
    call $parse_expr

    ;; Skip trailing whitespace
    local.get $end
    call $parse_skip_ws

    drop
)

;; Parse an expression
;; Returns expr, off
(func $parse_expr
    (param $off i32)
    (param $end i32)
    (result i32 i32)

    ;; Check for a left parenthesis
    local.get $off
    local.get $end
    call $parse_peek
    i32.const 0x28 ;; LEFT PARENTHESIS
    i32.eq
    (if (result i32)
        (then
            ;; There is a left parenthesis, this is a list/cell

            ;; Eat the left parenthesis
            local.get $off
            i32.const 1
            i32.add

            ;; Parse the rest of the list/cell
            local.get $end
            return_call $parse_list_inner
        )
        (else
            ;; There isn't a left parenthesis, this is an atom

            ;; Parse an atom
            local.get $off
            local.get $end
            return_call $parse_atom
        )
    )
)

;; Parse a list sans leading paren
;; Returns expr, off
(func $parse_list_inner
    (param $off i32)
    (param $end i32)
    (result i32 i32)
    (local $chr i32)

    ;; Skip leading whitespace
    local.get $off
    local.get $end
    call $parse_skip_ws

    ;; Check for a dot
    local.get $end
    call $parse_peek
    local.set $chr
    i32.const 0x2e ;; FULL STOP
    i32.eq
    (if (param i32) (result i32)
        (then
            ;; There is a dot, parse one final tail expression

            ;; Eat the dot
            i32.const 1
            i32.add

            ;; Skip whitespace after the dot
            local.get $end
            call $parse_skip_ws

            ;; Parse an expression
            local.get $end
            call $parse_expr

            ;; Skip trailing whitespace before the closing parenthesis
            local.get $end
            call $parse_skip_ws

            ;; Eat the closing paren
            i32.const 1
            i32.add
            return
        )
    )

    ;; Check for a right parenthesis
    local.get $chr
    i32.const 0x29 ;; RIGHT PARENTHESIS

```

```

i32.eq
(if (param i32) (result i32)
  (then
    ;; There is a right parenthesis, return the 0 atom

    local.set $off

    ;; Load the value zero
    i32.const 0

    ;; Eat the right parenthesis
    local.get $off
    i32.const 1
    i32.add
    return
  )
)

;; There are more items remaining

;; Parse an expression
local.get $end
call $parse_expr

;; Parse the rest of the list
local.get $end
call $parse_list_inner

;; Combine the results
local.set $off
call $cons
local.get $off
)

;; Parse an atom
;; Returns expr, off
(func $parse_atom
  (param $off i32)
  (param $end i32)
  (result i32 i32)
  (local $chr i32)
  (local $nex i32)

  ;; Determine the extent of the atom
  local.get $off
  local.get $off
  (block $break_loop (param i32) (result i32)
    (loop $loop (param i32) (result i32)
      local.get $end
      call $parse_peek
      local.tee $chr

      ;; Break for whitespace or end
      i32.const 32
      i32.le_u
      br_if $break_loop

      ;; Break for lparen
      local.get $chr
      i32.const 0x28
      i32.eq
      br_if $break_loop

      ;; Break for rparen
      local.get $chr
      i32.const 0x29
      i32.eq
      br_if $break_loop

      ;; Break for dot
      local.get $chr
      i32.const 0x2e
      i32.eq
      br_if $break_loop

      i32.const 1
      i32.add
      br $loop
    )
  )

  ;; Inter the atom
  local.tee $nex
  local.get $off
  i32.sub
  call $inter_string

  local.get $nex
)

;; Skip whitespace
(func $parse_skip_ws
  (param $off i32)
  (param $end i32)
  (result i32)

  ;; Loop over bytes
  local.get $off
  (block $break_loop (param i32) (result i32)
    (loop $loop (param i32) (result i32)
      local.get $end
      call $parse_peek

      ;; If the value is in [1, 32] it is whitespace
      i32.const 1
      i32.sub
      i32.const 32
      i32.ge_u
      br_if $break_loop
    )
  )

  ;; Enter the atom
  local.tee $nex
  local.get $off
  i32.sub
  call $inter_string

  local.get $nex
)

;; Parse an expression
local.get $end
call $parse_expr

;; Parse the rest of the list
local.get $end
call $parse_list_inner

;; Combine the results
local.set $off
call $cons
local.get $off
)

;; Table of builtins and system operations
;; The former take slots 0 - 16 while the latter start in slot 32
(table 64 funcref)
(elem (i32.const 0)
  $eval_builtin_quote
  $eval_builtin_true
  $eval_builtin_false
  $eval_builtin_head
  $eval_builtin_tail
  $eval_builtin_cons
  $eval_builtin_lte
  $eval_builtin_eq
  $eval_builtin_add
  $eval_builtin_sub
  $eval_builtin_and
  $eval_builtin_or
  $eval_builtin_not
  $eval_builtin_shl
  $eval_builtin_shr
  $eval_builtin_env
  $eval_builtin_sys
)

(type $invokable (func
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
))

;; Evaluate an expression in an environment, then perform garbage
;; collection from the given anchor.
;;
;; A garbage collection anchor is passed to this function so that the
;; needed garbage collection information can be maintained over tail calls.
;;
;; For the end user, you should probably just pass gc_get_anchor() unless
;; you would otherwise call gc_collect(<result of eval>, <some anchor>)
;; immediately after anyways.
(func $eval (export "eval")
  (param $expression i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  (local $receiver i32)

  ;; Determine whether the expression is an atom or cons cell
  local.get $expression
  i32.const 0
  i32.ge_s
  (if (result i32)
    (then
      ;; This is an atom, determine whether it is 0
      local.get $expression
      i32.eqz
      (if (result i32)
        (then
          ;; The 0 atom evaluates to itself
          i32.const 0
        )
        (else
          ;; Positive atoms are looked up in the environment
          local.get $expression
          local.get $environment
          call $lookup
        )
      )
    )
    (else
      ;; This is a cons cell, it should be a list representing an
      ;; invocation of either a user defined receiver or a builtin

      ;; Run garbage collection
      local.get $gc_anchor
      call $gc_collect

      ;; Evaluate the head to determine the receiver
    )
  )
)

```



```

    (param $arguments i32)
    (param $environment i32)
    (param $gc_anchor i32)
    (result i32)
    ;; Get the first argument
    local.get $arguments
    call $head

    ;; Tail evaluate it
    local.get $environment
    local.get $gc_anchor
    return_call $eval
)

;; Evaluate an invocation of the ~false builtin
(func $eval_builtin_false
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Tail evaluate it
  local.get $environment
  local.get $gc_anchor
  return_call $eval
)

;; Evaluate an invocation of the ~head builtin
(func $eval_builtin_head
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get its head
  call $head

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~tail builtin
(func $eval_builtin_tail
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get its tail
  call $tail

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~cons builtin
(func $eval_builtin_cons
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Cons them together
  call $cons

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~lte builtin
(func $eval_builtin_lte
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Prepare the ~true and ~false atoms
  i32.const 0x20_00_00_08 ;; ~true
  i32.const 0x20_00_00_10 ;; ~false

  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Select the appropriate atom (~true or ~false)
  ;; based on whether the first is less than or equal to the second
  i32.le_s
  select

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~eq builtin
(func $eval_builtin_eq
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Prepare the ~true and ~false atoms
  i32.const 0x20_00_00_08 ;; ~true
  i32.const 0x20_00_00_10 ;; ~false

  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Select the appropriate atom (~true or ~false)
  ;; based on whether the first is equal to the second
  i32.eq
  select

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~add builtin
(func $eval_builtin_add
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Add the arguments and take the result modulo 2 ^ 31
  i32.add
  i32.const 0x7F_FF_FF_FF
  i32.and

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

```

```

)    call $gc_collect

;; Evaluate an invocation of the ~sub builtin
(func $eval_builtin_sub
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Subtract the arguments and take the result modulo 2 ^ 31
  i32.sub
  i32.const 0x7F_FF_FF_FF
  i32.and

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~and builtin
(func $eval_builtin_and
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Take the bitwise and of the arguments
  i32.and

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~or builtin
(func $eval_builtin_or
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Take the bitwise or of the arguments
  i32.or

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~not builtin
(func $eval_builtin_not
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Take the bitwise negation of the least significant 31 bits
  i32.const 0x7F_FF_FF_FF
  i32.xor

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~shl builtin
(func $eval_builtin_shl
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Bit shift left the first argument by the second, keeping only the
  ;; least significant 31 bits
  i32.shl
  i32.const 0x7F_FF_FF_FF
  i32.and

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~shr builtin
(func $eval_builtin_shr
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the first argument
  local.get $arguments
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Get the second argument
  local.get $arguments
  call $tail
  call $head

  ;; Evaluate it
  local.get $environment
  global.get $cons_cells_top
  call $eval

  ;; Bit shift right the first argument by the second
  i32.shr_u

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~env builtin
(func $eval_builtin_env
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  ;; Get the environment
  local.get $environment

  ;; Run garbage collection
  local.get $gc_anchor
  call $gc_collect
)

;; Evaluate an invocation of the ~sys builtin
(func $eval_builtin_sys
  (param $arguments i32)
  (param $environment i32)
  (param $gc_anchor i32)
  (result i32)
  (local $code i32)

  ;; Get the argument

```



```

local.get $arguments
call $head

;; Evaluate it
local.get $environment
global.get $cons_cells_top
call $eval
local.tee $code

i32.eqz
(if (result i32)
  then
    call $construct_system_operation_table
  )
  else
    ;; Construct a system receiver
    ;; (this implementation represents these as (n . ~sys))
    local.get $code
    i32.const 0x20_00_00_80 ;; ~sys
    call $cons
  )
)

;; Run garbage collection
local.get $gc_anchor
call $gc_collect

)

;; ----- System Operation Registration -----

(type $system_operation_handler (func
  (param $arguments i32)
  (param $environment i32)
  (result i32)
))

(global $highest_system_opcode (mut i32) (i32.const 0))

;; Register a system operation under the designated name using the provided
;; handler. To be fully spec compliant this function should only be called
;; before any mu_ code has been evaluated.
(func (export "register_system_operation")
  (param $operation_name i32)
  (param $handler funcref)
  (local $opcode i32)
  (local $size_delta i32)

  ;; Get the next available opcode
  global.get $highest_system_opcode
  i32.const 1
  i32.add
  local.tee $opcode
  global.set $highest_system_opcode

  ;; Resize the function table as necessary

  i32.const 1
  i32.const 32

  ;; Compute the table index
  local.get $opcode
  i32.const 32
  i32.add

  ;; Round up to the next power of two
  i32.clz
  i32.sub
  i32.shl

  ;; Compute the size delta
  table.size
  i32.sub
  local.tee $size_delta

  ;; Resize if the delta is positive
  i32.const 0
  i32.gt_s
  (if
    then
      ref.null func
      local.get $size_delta
      table.grow
      drop
    )
  )

  ;; Set the designated slot
  local.get $opcode
  i32.const 32
  i32.add
  local.get $handler
  table.set

  ;; Update the string internment entry to indicate the opcode

  ;; Get the index into the internment stack from the atom number
  local.get $operation_name
  i32.const 0xDF_FF_FF_FF
  i32.and

  ;; Offset to the opcode entry
  i32.const 6
  i32.add

  ;; Write the opcode
  local.get $opcode
  i32.store16 (memory $string_internment_stack)
)

;; (~sys ()) -- get mappings from system operation names to codes

(func $construct_system_operation_table
  (result i32)
  (local $idx i32)
  (local $acc i32)
  (local $opcode i32)

  i32.const 0
  local.set $acc

  i32.const 0
  local.set $idx
  (loop $loop
    ;; Load the opcode
    local.get $idx
    i32.const 6
    i32.add
    i32.load16_u (memory $string_internment_stack)
    local.tee $opcode

    (block $skip_add (param i32)
      ;; If opcode is zero, this isn't a system operation, continue
      i32.eqz
      br_if $skip_add

      ;; Flip bit 29 of the index to get the atom number
      local.get $idx
      i32.const 0x20_00_00_00
      i32.xor

      ;; Add mapping from name to opcode to accumulator
      local.get $opcode
      call $cons
      local.get $acc
      call $cons
      local.set $acc
    )

    ;; Increment idx
    local.get $idx
    i32.const 8
    i32.add
    local.set $idx

    local.get $idx
    global.get $string_internment_stack_top
    i32.lt_u
    br_if $loop
  )

  local.get $acc
)

;; ----- Garbage Collector -----
;; The garbage collection mechanism operates as follows:
;; 1. Acquire a garbage collection anchor (gc_get_anchor)
;; 2. Do something that might cause the allocation of garbage
;; 3. Call gc_collect(<preserve>, <anchor>) passing the value you want to
;;    keep as preserve. The response is a new value (earlier on the stack
;;    where possible) which is equivalent to the passed preserve value.
;;    All cons cells which are not directly or indirectly a dependency of
;;    the preserved value and which were allocated after the anchor was
;;    taken are removed.
;;
;; Note that evaluation performs a garbage collection step itself so the
;; manual use of this mechanism by the embedder is only necessary to clear
;; cells created through a means other than evaluation such as parsing.

;; Get a garbage collection anchor
(func $gc_get_anchor (export "gc_get_anchor")
  (result i32)

  global.get $cons_cells_top
)

;; Run the garbage collector
(func $gc_collect (export "gc_collect")
  (param $preserve i32)
  (param $anchor i32)
  (result i32)
  (local $anchor_2 i32)

  ;; Take a second anchor, this represents where the cons stack grew to
  global.get $cons_cells_top
  local.set $anchor_2

  ;; Recursively copy the preserved element
  local.get $preserve
  local.get $anchor
  local.get $anchor_2
  call $gc_copy

  ;; Move the copied cells (the ones to be kept) down, overwriting the
  ;; range between the first and second anchors
  local.get $anchor
  i32.const 8
  i32.add
  local.get $anchor_2
  i32.const 8
  i32.add
  global.get $cons_cells_top
  local.get $anchor_2
  i32.sub
  memory.copy (memory $cons_cells)

  ;; Adjust the top of the stack down
  global.get $cons_cells_top
  local.get $anchor_2
  i32.sub
  local.get $anchor
  i32.add
)

```



```

    }

    match(value, pattern, environment = 0) {
        return this.#bindings.match(
            this.#conv(value),
            this.#conv(pattern),
            this.#conv(environment)
        );
    }

    eval(expression, environment = 0, anchor = null) {
        anchor ??= this.#bindings.gc_get_anchor();
        return this.#bindings.eval(
            this.#conv(expression),
            this.#conv(environment),
            anchor
        );
    }

    gc_get_anchor() {
        return this.#bindings.gc_get_anchor();
    }

    gc_collect(preserve, anchor) {
        return this.#bindings.gc_collect(this.#conv(preserve), anchor);
    }

    parse(str) {
        const { off, len } = this.#syalloc_string(str);
        return this.#bindings.parse(off, len);
    }

    // Method for rendering a mu_ object (which will just be an i32) as a
    // readable string. Implemented directly in Javascript as it is not part of
    // the core interpreter.
    //
    // The method chooses to render atoms without corresponding names as their
    // numeric representation as a u31 prefixed by the unicode symbol №.
    // Since the specification does not say anything about the interpretation
    // of characters outside the ascii range, this has no chance of colliding
    // with any code which is fully spec-compliant.
    //
    // This method only uses the dot symbol when it has to, always preferring
    // to render lists as much as it can.
    show(obj) {
        obj = this.#conv(obj);
        if (obj > 0) {
            const [ off, len ] = this.#bindings.lookup_interred_string(obj);
            if (len === -1) {
                return '№' + obj.toString();
            } else {
                const buf = new Uint8Array(
                    this.#bindings.stringyard.buffer,
                    off,
                    len
                );
                return this.#text_decoder.decode(buf);
            }
        } else {
            let str = '(';
            let first = true;
            while (obj < 0) {
                if (first) first = false;
                else str += ' ';
                str += this.show(this.head(obj));
                obj = this.tail(obj);
            }

            if (obj > 0) {
                str += " . ";
                str += this.show(obj);
            }

            str += ')';
            return str;
        }
    }

    // Convert a Javascript object into a mu_ one as best as possible
    // Converts arrays to lists and atoms to their atom numbers.
    #conv(obj) {
        if (typeof obj === "number") {
            return obj;
        } else if (typeof obj === "string") {
            return this.parse(obj);
        } else if (Array.isArray(obj)) {
            return obj.reduceRight(
                (acc, val) => this.cons(val, acc),
                this.#conv(obj["tail"] ?? 0)
            );
        } else if (obj === null) {
            return 0;
        }
    }

    // Inter a Javascript string into the mu_ interpreter
    #inter_string(str) {
        const { off, len } = this.#syalloc_string(str);
        return this.#bindings.inter_string(off, len);
    }

    #text_encoder = new TextEncoder();
    #text_decoder = new TextDecoder();

    // Allocate a javascript string into the mu_ interpreter's space
    #syalloc_string(str) {
        // We ask for 3 bytes per UTF-16 unit since that is an upper bound
        const size = str.length * 3;
        const off = this.#bindings.syalloc(size);
        const buf = new Uint8Array(
            this.#bindings.stringyard.buffer,
            off,
            size
        );
        const { written: len } = this.#text_encoder.encodeInto(str, buf);
        return { off, len };
    }
}

```

Appendix E

Tests